

Homework 5: Data Platforms for Data Engineering Zoomcamp 2026

<https://github.com/DataTalksC...>

Due date: 28 February 2026 20:59 (local time)

Questions

Question 1. Bruin Pipeline Structure In a Bruin project, what are the required files/directories? (1 point)

- ☐ bruin.yml and assets/
- ☐ .bruin.yml and pipeline.yml (assets can be anywhere)
- ☒ .bruin.yml and pipeline/ with pipeline.yml and assets/
- ☐ pipeline.yml and assets/ only

Question 2. Materialization Strategies You're building a pipeline that processes NYC taxi data organized by month based on pickup_datetime. Which incremental strategy is best for processing a specific interval period by deleting and inserting data for that time period? (1 point)

- ☐ append - always add new rows
- ☐ replace - truncate and rebuild entirely
- ☒ time_interval - incremental based on a time column
- ☐ view - create a virtual table only

Question 3. Pipeline Variables You have a variable defined in pipeline.yml: variables: taxi_types: type: array items: type: string default: ["yellow", "green"] How do you override this when running the pipeline to only process yellow taxis? (1 point)

- ☐ bruin run --taxi-types yellow
- ☐ bruin run --var taxi_types=yellow
- ☒ bruin run --var 'taxi_types=["yellow"]'
- ☐ bruin run --set taxi_types=["yellow"]

Question 4. Running with Dependencies You've modified the ingestion/trips.py asset and want to run it plus all downstream assets. Which command should you use? (1 point)

- ☐ bruin run ingestion.trips --all
- ☒ bruin run ingestion/trips.py --downstream
- ☐ bruin run pipeline/trips.py --recursive
- ☐ bruin run --select ingestion.trips+

Question 5. Quality Checks. You want to ensure the pickup_datetime column in your trips table never has NULL values. Which quality check should you add to your asset definition? (1 point)

- ☐ name: unique
- ☒ name: not_null
- ☐ name: positive
- ☐ name: accepted_values, value: [not_null]

Question 6. Lineage and Dependencies After building your pipeline, you want to visualize the dependency graph between assets. Which Bruin command should you use? (1 point)

- ☐ bruin graph
- ☐ bruin dependencies
- ☒ bruin lineage
- ☐ bruin show

Question 7. First-Time Run You're running a Bruin pipeline for the first time on a new DuckDB database. What flag should you use to ensure tables are created from scratch? (1 point)

- ☐ --create
- ☐ --init
- ☒ --full-refresh
- ☐ --truncate

Homework URL

<https://github.com/dammi01/data-engineering-zoomcamp>

Learning in public links (optional)

<https://x.com/mdambock/status/2027113466390007981?s=20>

https://www.linkedin.com/posts/michael-dambock_dataengineeringzoomcamp-dataengineeri

Add Link

Time spent on lectures (hours) (optional)

4

Time spent on homework (hours) (optional)

6

FAQ contribution (FAQ document, optional)

Memory errors. Restricting the fetch to only necessary columns (VendorID, dates, amounts) significantly improves performance without losing analytical value for the homework.

Submit

Status: Not submitted